

NAME: OMAR USMAN UPPAL

CMS ID : 203515

ASSIGNMENT # 3

BIG DATA : HADOOP – HIVE – HBASE - SPARK

1. QUESTION 1 – HADOOP

HADOOP INTALLATION:

Following images show different steps involved in HADOOP Installation

Image 1: Hadoop 2.7.5 was downloaded and installed on a UBUNTU Virtual Machine with a RAM of 3 GB and 40 GB Hard Disk space.

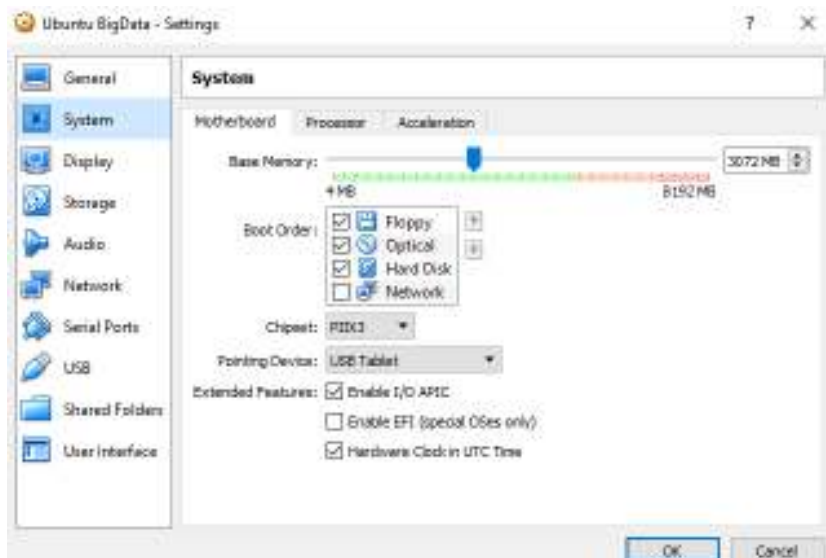


Image 2

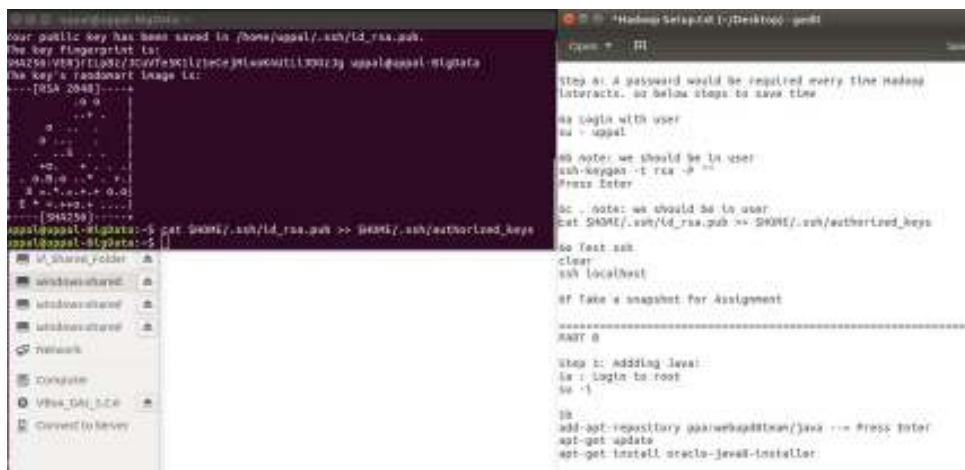


Image 3



Image 4

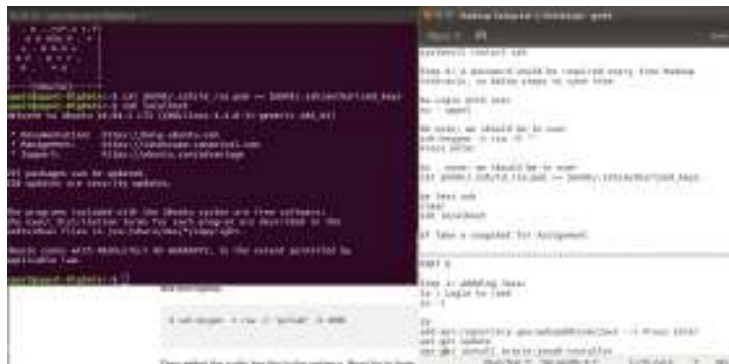


Image 5

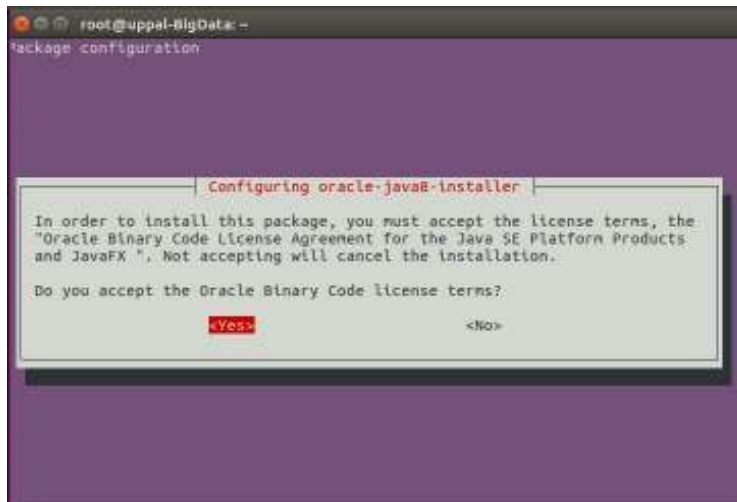


Image 6




Image 7

Image 8

Image 9

Image 10

HADOOP COMMANDS:

Command 1: `$ hadoop fs -ls -R /`

This command shows the complete directory system of the Hadoop file system including the files inside the folder.

```

uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -ls -R /
drwxr-xr-x - uppal supergroup 0 2018-04-12 06:01 /hbase
drwxr-xr-x - uppal supergroup 0 2018-04-12 06:03 /hbase/.tmp
drwxr-xr-x - uppal supergroup 0 2018-04-12 06:03 /hbase/.tmp/data
drwxr-xr-x - uppal supergroup 0 2018-04-12 06:04 /hbase/.tmp/data/default
drwxr-xr-x - uppal supergroup 0 2018-04-12 06:09 /hbase/MasterProcWALS
-rw-r--r-- 1 uppal supergroup 30 2018-04-12 13:47 /hbase/MasterProcWALS/state-00000000000000000015.log
drwxr-xr-x - uppal supergroup 0 2018-04-12 06:01 /hbase/WALS

```

Command 2: \$ hadoop fs -mkdir /user/uppal/Hadoop

“hadoop fs -mkdir” is used to create a folder in the Hadoop file system. The datasets downloaded are inside Ubuntu file system and they need to be accessed by Hadoop file system so Hadoop commands could be executed on those files. Therefore, a new folder is created inside another Hadoop accessed directory.

```

uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -mkdir /user/uppal/hadoop

```

Command 3: \$ hadoop fs -put /tmp/Q1 /user/uppal/hadoop

3 datasets were downloaded from website: <https://www.michael-noll.com/tutorials/writing-an-hadoop-mapreduce-program-in-python/> and renamed as “doc1”, “doc2” and “doc3”. Using “put” command datasets were moved to Hadoop folder “/user/uppal/hadoop/Q1”.

Command 4: \$ hadoop fs -ls -R /user/uppal/Hadoop

Datasets were stored in ../Q1 folder and the folder was moved to Hadoop folder ../Hadoop. It can be seen in the below snapshots.

```

uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -ls -R /user/uppal/hadoop
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -put /tmp/Q1 /user/uppal/hadoop
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -ls -R /user/uppal/hadoop
drwxr-xr-x - uppal supergroup 0 2018-04-14 13:49 /user/uppal/hadoop/Q1
-rw-r--r-- 4 uppal supergroup 1495 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc1.txt
-rw-r--r-- 4 uppal supergroup 5156 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc2.txt
-rw-r--r-- 4 uppal supergroup 2200 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc3.txt
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$

```

Command 5: \$ hadoop fs -cat /user/uppal/hadoop/doc3.txt

“CAT” command was used to print the content of the dataset on the console.


```
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -cat /user/uppal/hadoop/Q1/doc3.txt
```

The Project Gutenberg EBook of Ulysses, by James Joyce

This eBook is for the use of anyone anywhere at no cost and with almost no restrictions whatsoever. You may copy it, give it away or re-use it under the terms of the Project Gutenberg License included with this eBook or online at www.gutenberg.org

Mr Bloom stood at the corner, his eyes wandering over the multicoloured hoardings. Cantrell and Cochrane's Ginger Ale (Aromatic). Clery's Summer Sale. No, he's going on straight. Hello. Leah tonight. Mrs Bandmann Palmer. Like to see her again in that. Hamlet she played last night. Male impersonator. Perhaps he was a woman. Why Ophelia committed suicide. Poor papa! How he used to talk of Kate Bateman in that. Outside the Adelphi in London waited all the afternoon to get in. Year before I was born that was: sixtyfive. And Ristori in Vienna. What is this the right name is? By Mosenthal it is. Rachel, is it? No. The scene he was always talking about where the old blind Abraham recognises the voice and puts his fingers on his face.

He stood aside watching their blind masks pass down the aisle, one by one, and seek their places. He approached a bench and seated himself in its corner, nursing his hat and newspaper. These pots we have to wear. We ought to have hats modelled on our heads. They were about him here and there, with heads still bowed in their crimson halters, waiting for it to melt in their stomachs. Something like those mazzoth: it's that sort of bread: unleavened shewbread. Look at them. Now I bet it makes them feel happy. Lollipop. It does. Yes, bread of angels it's called. There's a big idea behind it, kind of kingdom of God is within you feel. First communicants. Hokypoky penny a lump. Then feel all like one family party, same in the theatre, all in the same swim. They do. I'm sure of that. Not so lonely. In our confraternity. Then come out a bit spreeish. Let off steam. Thing is if you really believe in it. Lourdes cure, waters of oblivion, and the Knock apparition, statues bleeding. Old fellow asleep near that confessionbox. Hence those snores. Blind faith. Safe in the arms of kingdom come. Lulls all pain. Wake this time

Command 6: \$ hadoop fs -rm -r /user/uppal/hadoop/Q2

Q2 folder was created and then using “-rm -r” command the folder was removed from the Hadoop file system as it can be seen in below set of snapshots.

```
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -mkdir /user/uppal/hadoop/Q2
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -ls -R /user/uppal/hadoop
drwxr-xr-x  - uppal supergroup          0 2018-04-14 13:49 /user/uppal/hadoop/Q1
-rw-r--r--  4 uppal supergroup       1495 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc1.txt
-rw-r--r--  4 uppal supergroup       5156 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc2.txt
-rw-r--r--  4 uppal supergroup       2200 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc3.txt
drwxr-xr-x  - uppal supergroup          0 2018-04-14 14:46 /user/uppal/hadoop/Q2
```

```
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -rm -r /user/uppal/hadoop/Q2
```

```
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -mkdir /user/uppal/hadoop/Q2
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -rm -r /user/uppal/hadoop/Q2
18/04/14 14:49:04 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Empty interval = 0 minutes.
Deleted /user/uppal/hadoop/Q2
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -ls -R /user/uppal/hadoop
drwxr-xr-x  - uppal supergroup          0 2018-04-14 13:49 /user/uppal/hadoop/Q1
-rw-r--r--  4 uppal supergroup       1495 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc1.txt
-rw-r--r--  4 uppal supergroup       5156 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc2.txt
-rw-r--r--  4 uppal supergroup       2200 2018-04-14 13:49 /user/uppal/hadoop/Q1/doc3.txt
uppal@uppal-BigData:/usr/local/hadoop/hadoop-2.7.5/sbin$
```

MAP-REDUCE function using Hadoop Streaming:

Hadoop streaming is a utility that comes with the Hadoop distribution. The utility allows you to create and run Map/Reduce jobs with any executable or script as the mapper and/or the reducer. The mapper and reducer codes are written in python as below:

MAPPER CODE:

Mapper code will read the data from STDIN and splits into key-value pairs (word, count=1). The output is list of lines of the pairs to STDOUT.

```
#!/usr/bin/env python
"""mapper.py"""

import sys

# input comes from STDIN (standard input)
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into words
    words = line.split()
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        #
        # tab-delimited; the trivial word count is 1
        print '%s\t%s' % (word, 1)
```

REDUCER CODE:

Reducer code will read the results of mapper function from STDIN and sums the occurrence of each word to final count. The output is list of lines of the pairs (words, sum) to STDOUT.

```
#!/usr/bin/env python
"""reducer.py"""

from operator import itemgetter
import sys

current_word = None
current_count = 0
word = None

# input comes from STDIN
for line in sys.stdin:
    # remove leading and trailing whitespace

    line = line.strip()

    # parse the input we got from mapper.py
    word, count = line.split('\t', 1)

    # this IF-switch only works because Hadoop sorts map output
    # by key (here: word) before it is passed to the reducer
    if current_word == word:
        current_count += count
    else:
        if current_word:
            # write result to STDOUT
            print '%s\t%s' % (current_word, current_count)
            current_count = count
            current_word = word
# do not forget to output the last word if needed!
if current_word == word:
    print '%s\t%s' % (current_word, current_count)
```

Command 7: \$ chmod +x /usr/local/hadoop/mapper.py

The command is run for getting the execution permission of the mapper.py file else it will run into problems. Same permission is give reducer.py.

STRUCTURE OF HADOOP STREAMING COMMAND:

mapred streaming --> streaming command *.jar

-input myInputDirs --> directory containing input files

-output myOutputDir --> output directory will be created. It should not exist else it will give error

-mapper PythonMap.py --> mapper script

-reducer PythonReduce.py --> reducer script

Command 8:

**hadoop jar ./hadoop-2.7.5/share/hadoop/tools/lib/hadoop-streaming-2.7.5.jar **

**-file /usr/local/hadoop/mapper.py -mapper /usr/local/hadoop/mapper.py **

**-file /usr/local/hadoop/reducer.py -reducer /usr/local/hadoop/reducer.py **

-input /user/uppal/hadoop/Q1/* -output /user/uppal/hadoop/Q1-output

```
uppal@uppal-BigData:/usr/local/hadoop$ hadoop jar ./hadoop-2.7.5/share/hadoop/tools/lib/hadoop-streaming-2.7.5.jar -file /usr/local/hadoop/mapper.py -mapper /usr/local/hadoop/mapper.py -file /usr/local/hadoop/reducer.py -reducer /usr/local/hadoop/reducer.py -input /user/uppal/hadoop/Q1/* -output /user/uppal/hadoop/Q1-output
```

OUTPUT:

Output of the command was stored in a newly formed output directory as give in the image. The name of the output file of reducer is “part-00000”

```
uppal@uppal-BigData:/usr/local/hadoop$ hadoop fs -ls -R /user/uppal/hadoop/Q1-output
-rw-r--r--  4 uppal supergroup      0 2018-04-15 00:31 /user/uppal/hadoop/Q1-output/_SUCCESS
-rw-r--r--  4 uppal supergroup  7351 2018-04-15 00:31 /user/uppal/hadoop/Q1-output/part-00000
uppal@uppal-BigData:/usr/local/hadoop$ hadoop fs -cat /user/uppal/Q1-output/part-00000
cat: '/user/uppal/Q1-output/part-00000': No such file or directory
uppal@uppal-BigData:/usr/local/hadoop$ hadoop fs -cat /user/uppal/hadoop/Q1-output/part-00000
```

Executing the “cat” command displays a long list of key-value pairs. A small part is shown in the image below:

"Outline		1
"he	1	
"legal	1	
"primers,"		1
"that	1	
(#3	1	
(Aromatic).		1
1922	10	
40	1	
800	1	
A	3	
ABERDEEN		1
ABOUT	1	
AND	1	
ARE	1	
ARTHUR	1	
Abraham	1	
Adelphi	1	
Ale	1	
America	1	

2. QUESTION 2 - HIVE

Download Dataset

Bird.csv dataset is downloaded from the link:

<http://stat-computing.org/dataexpo/2011/resources/data/birds.csv>

Copy the file to Hadoop folder

```
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -copyFromLocal /tmp/birds.csv /tmp/hive/birds.csv
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -ls -R /tmp/hive
-rw-r--r--  4 uppal supergroup    635111 2018-04-13 17:55 /tmp/hive/birds.csv
drwx-----  - uppal supergroup      0 2018-04-13 17:52 /tmp/hive/uppal
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$

-rw-r--r--  1 uppal supergroup    42 2018-04-06 19:32 /hbase/hbase.id
-rw-r--r--  1 uppal supergroup    7 2018-04-06 19:32 /hbase/hbase.version
drwxr-xr-x  - uppal supergroup    0 2018-04-12 06:12 /hbase/oldWALS
-rw-r--r--  4 uppal supergroup   280 2018-04-09 10:36 /lnpt
drwx-wx-wx  - uppal supergroup    0 2018-04-09 10:44 /tmp
drwx-----  - uppal supergroup    0 2018-04-09 10:44 /tmp/hadoop-yarn
drwx-----  - uppal supergroup    0 2018-04-09 10:44 /tmp/hadoop-yarn/staging
drwxr-xr-x  - uppal supergroup    0 2018-04-09 10:44 /tmp/hadoop-yarn/staging/history
drwxrwxrwt  - uppal supergroup    0 2018-04-09 10:44 /tmp/hadoop-yarn/staging/history/done_intermediate
drwxrwxrwt  - uppal supergroup    0 2018-04-14 10:04 /tmp/hadoop-yarn/staging/history/done_intermediate/uppal
-rwxrwxr-x  4 uppal supergroup   23326 2018-04-09 10:45 /tmp/hadoop-yarn/staging/history/done_intermediate/uppal/job_1523251470863_0001-15
23252687957-uppal-streamjob737540380446225538.jar-1523252707607-0-1-SUCCEEDED-default-1523252699372.jhist
-rwxrwxr-x  4 uppal supergroup    355 2018-04-09 10:45 /tmp/hadoop-yarn/staging/history/done_intermediate/uppal/job_1523251470863_0001.su
```

Operation 1: Starts hive shell → \$ hive

Operation 2: Create Table → \$ create table Birds_table(Species STRING, Latitude FLOAT, Longitude FLOAT, Oiling STRING, Condition STRING, BirdCount INT, Date_1 STRING, Oil_Cond INT, Date_2 STRING, week_number INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE LOCATION '/user/uppal/hive';

The command will create a table “Birds_table” with column names give in the brakcets with their data type. For Example, “Species” is a column name with “STRING” as data type. Since it is a *.csv file so delimitation criteria is give and then table is stored at a specific location

Operation 3: Check table created? → SHOW TABLES;

```
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$ hive
Logging initialized using configuration in jar:file:/usr/lib/hive/apache-hive-1.2.2-bin/lib/hive-common-1.2.2.jar!/hive-log4j.properties
hive> create table Birds_table(Species STRING, Latitude FLOAT, Longitude FLOAT, Oiling STRING, Condition STRING, BirdCount INT, Date_1 STRING,
Oil_Cond INT, Date_2 STRING, week_number INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE LOCATION '/user/uppal/hive';
OK
Time taken: 1.598 seconds
hive> SHOW TABLES;
OK
birds_table
names2
names_text
Time taken: 0.241 seconds, Fetched: 3 row(s)
```

The command shows the tables available in the HIVE.

Operation 4: Check table information: → DESCRIBE “Birds_table”

We can see that this command shows the column names and their data types of the table. It also shows the count that there are 10 columns and the time it took to provide the information.

```
hive> DESCRIBE birds_table;
OK
species                string
latitude               float
longitude              float
oiling                 string
condition              string
birdcount              int
date_1                 string
oil_cond               int
date_2                 string
week_number            int
Time taken: 0.2 seconds, Fetched: 10 row(s)
hive>
```

Operation 5: LOAD Data: LOAD DATA INPATH '/tmp/hive/birds.csv' OVERWRITE INTO TABLE Birds_table;

The data from the dataset “birds.csv” will be uploaded in the table. Once done we can execute query commands on the table.

```
hive> LOAD DATA INPATH '/tmp/hive/birds.csv' OVERWRITE INTO TABLE Birds_table;
Loading data to table default.birds_table
Table default.birds_table stats: [numFiles=0, numRows=0, totalSize=0, rawDataSize=0]
OK
Time taken: 1.02 seconds
hive>
```

Query 1: SELECT * from birds_table limit 5;

The query returns 5 arbitrary birds data.

```
hive> SELECT * from birds_table limit 5;
OK
"Species"      NULL      NULL      "Oiling"      "Condition"      NULL      "Date"      NULL      "Date"      NULL
"Northern Gannet" 30.3286 -89.1981 "Not Visibly Oiled" "Live" 1      2010-07-21 1      2010-07-21 30
"Laughing Gull" 30.23172 -88.32127 "Not Visibly Oiled" "Live" 1      2010-05-05 1      2010-05-05 19
"Northern Gannet" 30.26677 -87.59248 "Visibly Oiled" "Live" 1      2010-05-05 2      2010-05-05 19
"American White Pelican" 29.29649 -89.66432 "Not Visibly Oiled" "Live" 1      2010-05-05 1      2010-05-05 1
9
Time taken: 0.435 seconds, Fetched: 5 row(s)
hive> alter table birds_table set tblproperties ("skip.header.line.count"="1");
OK
Time taken: 0.262 seconds
hive> SELECT * from birds_table limit 5;
OK
"Species"      NULL      NULL      "Oiling"      "Condition"      NULL      "Date"      NULL      "Date"      NULL
"Northern Gannet" 30.3286 -89.1981 "Not Visibly Oiled" "Live" 1      2010-07-21 1      2010-07-21 30
"Laughing Gull" 30.23172 -88.32127 "Not Visibly Oiled" "Live" 1      2010-05-05 1      2010-05-05 19
"Northern Gannet" 30.26677 -87.59248 "Visibly Oiled" "Live" 1      2010-05-05 2      2010-05-05 19
"American White Pelican" 29.29649 -89.66432 "Not Visibly Oiled" "Live" 1      2010-05-05 1      2010-05-05 1
9
"Brown Pelican" 29.08244 -88.07624 "Visibly Oiled" "Live" 1      2010-05-08 2      2010-05-08 19
Time taken: 0.209 seconds, Fetched: 5 row(s)
hive>
```

Query 2: SELECT species FROM birds_table WHERE Oil_cond =1;

The query returns list of Species where oiling condition is 1. There are 814 rows given at the end of the result while the small subset of output format is shown in the figure.

```
hive> SELECT a.species FROM birds_table a WHERE a.Oil_cond =1;
OK
"Northern Gannet"
"Laughing Gull"
"American White Pelican"
"Brown Pelican"
"Common Loon"
"Northern Gannet"
"Brown Pelican"
"Laughing Gull"
```

```
Royal Tern
"Magnificent Frigatebird"
"Laughing Gull"
"Laughing Gull"
Time taken: 0.437 seconds, Fetched: 814 row(s)
hive> █
```

Query 3: SELECT a.species FROM birds_table a WHERE a.Oil_cond =1 limit 10;

The query returns 10 arbitrary name of species that meet the criteria of Oil_cond=1;

```
hive> SELECT a.species FROM birds_table a WHERE a.Oil_cond =1 limit 10;
OK
"Northern Gannet"
"Laughing Gull"
"American White Pelican"
"Brown Pelican"
"Common Loon"
"Northern Gannet"
"Brown Pelican"
"Laughing Gull"
"Brown Pelican"
"Herring Gull"
Time taken: 0.229 seconds, Fetched: 10 row(s)
hive> █
```

Query 4: SELECT * FROM birds_table a WHERE a.week_number =31;

The query returns all the data of week.number 31. It shows there are 592 rows of such type.

```
hive> SELECT * FROM birds_table a WHERE a.week_number =31;
```

```
Time taken: 0.189 seconds, Fetched: 592 row(s)
```

Query 5: SELECT Count(*) FROM birds_table a WHERE a.week_number =31;

There are 592 samples from Week 31. This can be checked using Count(*).

```
hive> SELECT Count(*) FROM birds_table a WHERE a.week_number =31;
Query ID = uppal_20180413183538_8afdd0b2-0980-464e-b3d3-784714122c0b
Total jobs = 1
```

Query 6: SELECT DISTINCT SPECIES FROM birds_table WHERE week_number = 31 AND OIL_COND = 1;

Using the “DISTINCT” keyword we can retrieve the list of species without duplicate that were affected by oil_cond=1 in week.number 31.

```
hive> SELECT DISTINCT SPECIES FROM birds_table WHERE week_number = 31 AND OIL_COND = 1;
Query ID = uppal_20180414094816_a2b2a801-81ae-4af7-a4dc-d7a7fad4df21
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523627382222_0007, Tracking URL = http://uppal-BigData:8088/proxy/application_1523627382222_0007/
Kill Command = /usr/local/hadoop/hadoop-2.7.5/bin/hadoop job -kill job_1523627382222_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-04-14 09:40:28,929 Stage-1 map = 0%, reduce = 0%
2018-04-14 09:48:39,891 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.86 sec
2018-04-14 09:48:50,689 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.98 sec
MapReduce Total cumulative CPU time: 3 seconds 980 nsec
Ended Job = job_1523627382222_0007
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.98 sec HDFS Read: 643661 HDFS Write: 199 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 980 nsec
OK
"Black Crowned Night Heron"
"Black Skinner"
"Brown Pelican"
"Clapper Rail"
"Common Loon"
"Laughing Gull"
"Mallard"
"Northern Gannet"
"Other"
"Royal Tern"
"Unidentified Dowitcher"
"Unidentified Gull"
Time taken: 36.425 seconds, Fetched: 12 row(s)
hive>
```


3. QUESTION 3 - HBASE

Download Dataset

turtles.csv dataset is downloaded from the link:

<http://stat-computing.org/dataexpo/2011/resources/data/turtles.csv>

Operation 1: \$ jps

“jps” displays different daemons running in the background. We can see that “hbase” is not active.

```
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$ jps
4049 ResourceManager
3665 DataNode
4178 NodeManager
4228 Jps
3832 SecondaryNameNode
3544 NameNode
```

Operation 2: \$ start-hbase.sh

This command will start the hbase and we can see the result using “jps” command.

```
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$ jps
4049 ResourceManager
5217 Jps
3665 DataNode
4178 NodeManager
3832 SecondaryNameNode
3544 NameNode
4985 HRegionServer
4857 HMaster
4794 HQuorumPeer
```

Operation 3: \$ hbase shell

We will enter the hbase shell and with that we can create tables in the shell and start our query operations.

Operation 4: \$ version

The command will result the information about the “hbase”. It includes the version of Hbase and its release etc. This command will be executed inside the shell.

Operation 5: \$ whoami

It will give the information of the user currently active and its permission status regarding “Hbase”. This command will be executed inside the shell.

Operation 6 - Create table: create "turtles", "turtles_data"

The command will be executed inside hbase shell and table “turtles” will be created. “turtles_data” is used for columns family name.

Operation 7 - \$ list

It will display the available tables in the hbase system. Table is created and now will exit the hbase shell and execute “hbase” command to load the data in the table.

```
hbase(main):004:0> create "turtles" , "turtles_data"
0 row(s) in 1.8090 seconds

=> Hbase::Table - turtles
hbase(main):005:0> list
TABLE
birds
birds_2
new
turtles
4 row(s) in 0.0900 seconds
```

Operation 8 - Copy the file to Hadoop folder: \$ `hadoop fs -copyFromLocal /tmp/turtles.csv /user/uppal/hbase/turtles.csv`

```
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -copyFromLocal /tmp/turtles.csv /user/uppal/hbase/turtles.csv
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$ hadoop fs -ls -R /user/uppal/hbase/turtles.csv
-rw-r--r--  4 uppal supergroup    75938 2018-04-15 06:35 /user/uppal/hbase/turtles.csv
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$
```

Operation 9: Load the Data: \$ `hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.separator="," -Dimporttsv.columns="HBASE_ROW_KEY,turtles_data:obsData,turtles_data:State,turtles_data:Latitude,turtles_data:Longitude,turtles_data:New,turtles_data:Confirmed,turtles_data:Alive,turtles_data>Date,turtles_data:week.number" turtles /user/uppal/hbase/turtles.csv`

This command will load the data in the “turtles” table. Column name “HBASE_ROW_KEY” is used to designate that this column should be used as the row key for each imported record. In our case it is the “species” column of “turtles.csv”, which will now be the “row key”.

```
uppal@uppal-BigData: /usr/local/hadoop/hadoop-2.7.5/sbin$ hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.separator="," -Dimporttsv.columns="HBASE_ROW_KEY,birds_data:Latitude,birds_data:Longitude,birds_data:Oiling,birds_data:Condition,birds_data:BirdCount,birds_data:Date,birds_data:Oil_Cond,birds_data>Date,birds_data:week.number" birds /user/uppal/hbase/birds.csv
```

Operation 10- Hbase Shell again: \$ `Hbase shell`

Table is loaded with data and we will execute our queries inside the hbase shell.

Query 1: `get "turtles", "Caretta caretta"`

The query returns the values of all columns against row id of “Caretta caretta”.

```
hbase(main):002:0> get "turtles", "Caretta caretta"
COLUMN                                CELL
turtles_data:Alive                    timestamp=1523756326651, value=N
turtles_data:Confirmed                 timestamp=1523756326651, value=Y
turtles_data:Date                     timestamp=1523756326651, value=2010-09-27
turtles_data:Latitude                  timestamp=1523756326651, value=30.29833
turtles_data:Longitude                 timestamp=1523756326651, value=-87.55388
turtles_data:New                       timestamp=1523756326651, value=N
turtles_data:State                     timestamp=1523756326651, value=AL
turtles_data:obsData                   timestamp=1523756326651, value=2010-09-27
turtles_data:week.number                timestamp=1523756326651, value=40
9 row(s) in 0.1930 seconds
```

Query 2: get "turtles", "Caretta caretta", {COLUMN => "turtles_data:week.number"}

The query returns the values of all column “week.number” which is 40 against row id of “Caretta caretta”.

```
hbase(main):004:0> get "turtles", "Caretta caretta", {COLUMN => "turtles_data:week.number"}
COLUMN                                CELL
turtles_data:week.number                timestamp=1523756326651, value=40
1 row(s) in 0.1540 seconds
```

Query 3: get "turtles", "Caretta caretta", {COLUMN => "turtles_data:week.number", COLUMN => "turtles_data:Alive"}

The query returns the values of week.number and if it is alive or not against row value of specie “Caretta caretta”.

```
hbase(main):010:0> get "turtles", "Caretta caretta", {COLUMN => "turtles_data:week.number", COLUMN => "turtles_data:Alive"}
COLUMN                                CELL
turtles_data:Alive                    timestamp=1523756326651, value=N
1 row(s) in 0.0360 seconds
```

Query 4: scan "turtles", { COLUMNS => "turtles_data:week.number", LIMIT => 3}

This query scans the table “turtles” and returns the three rows. The value shown is of 1 column only which is “week.number”.

```
hbase(main):002:0> scan "turtles", { COLUMNS => "turtles_data:week.number", LIMIT => 3}
ROW                                COLUMN+CELL
Caretta caretta                    column=turtles_data:week.number, timestamp=1523756326651, value=40
Chelonia mydas                     column=turtles_data:week.number, timestamp=1523756326651, value=42
Eretmochelys imbricata              column=turtles_data:week.number, timestamp=1523756326651, value=38
3 row(s) in 0.0950 seconds
hbase(main):003:0> █
```

Query 5: scan "turtles", { FILTER => "KeyOnlyFilter()"}

This filter does not take any arguments. It returns only the key component of each k.

```
hbase(main):003:0> scan "turtles", { FILTER => "KeyOnlyFilter()" }
ROW                                COLUMN+CELL
Caretta caretta                   column=turtles_data:Alive, timestamp=1523756326651, value=
Caretta caretta                   column=turtles_data:Confirmed, timestamp=1523756326651, value=
Caretta caretta                   column=turtles_data:Date, timestamp=1523756326651, value=
Caretta caretta                   column=turtles_data:Latitude, timestamp=1523756326651, value=
Caretta caretta                   column=turtles_data:Longitude, timestamp=1523756326651, value=
Caretta caretta                   column=turtles_data:New, timestamp=1523756326651, value=
Caretta caretta                   column=turtles_data:State, timestamp=1523756326651, value=
Caretta caretta                   column=turtles_data:obsData, timestamp=1523756326651, value=
Caretta caretta                   column=turtles_data:week.number, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:Alive, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:Confirmed, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:Date, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:Latitude, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:Longitude, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:New, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:State, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:obsData, timestamp=1523756326651, value=
Chelonia mydas                    column=turtles_data:week.number, timestamp=1523756326651, value=
```

Query 6: get "turtles", "Lepidochelys kempii"

The query returns the values of all columns against row id of "Lepidochelys kempii".

```
hbase(main):014:0> get "turtles", "Lepidochelys kempii"
COLUMN                                CELL
turtles_data:Alive                  timestamp=1523756326651, value=N
turtles_data:Confirmed              timestamp=1523756326651, value=Y
turtles_data:Date                   timestamp=1523756326651, value=2010-10-18
turtles_data:Latitude               timestamp=1523756326651, value=30.231953
turtles_data:Longitude              timestamp=1523756326651, value=-88.317361
turtles_data:New                    timestamp=1523756326651, value=Y
turtles_data:State                  timestamp=1523756326651, value=AL
turtles_data:obsData                timestamp=1523756326651, value=2010-10-18
turtles_data:week.number            timestamp=1523756326651, value=43
9 row(s) in 0.1000 seconds

hbase(main):015:0>
```

Query 7: Delete Table:

To delete a table, first we have to disable it and then we can drop the table. It is shown in the image below:


```
hbase(main):010:0> list
TABLE
birds
birds_2
new
turtles
4 row(s) in 0.4010 seconds

=> ["birds", "birds_2", "new", "turtles"]
hbase(main):011:0> disable "birds_2"
0 row(s) in 2.9880 seconds

hbase(main):012:0> drop "birds_2"
0 row(s) in 2.4440 seconds

hbase(main):013:0> list
TABLE
birds
new
turtles
3 row(s) in 0.0180 seconds

=> ["birds", "new", "turtles"]
hbase(main):014:0> 
```


4. SPARK BASIC OPERATIONS

Basic Operations:

Some of the basic Operation of SPARK are below mentioned in the screenshots with their results.

Operation 1: `sc.version ()` returns the version of PySpark.

Operation 2: `sc.pythonVer` returns the verison of python.

Operation 3: `str(sc.sparkUser())` returns the information about the current Spark user.

Operation 4: `sc.appName` retruns the name of the app.

```
In [133]: sc.version
Out[133]: '2.2.0'

In [134]: sc.pythonVer
Out[134]: '3.6'

In [135]: str(sc.sparkUser())
Out[135]: 'uppal'

In [136]: sc.appName
Out[136]: 'PySparkShell'
```

Data Operations:

File can be read using Spark Context. To check the basic commands of Spark the dataset we are using is “cars.csv” downloaded from link “<http://lib.stat.cmu.edu/datasets/cars.data>”. Preprocessing of the data is not done but other than that we can apply some basic operations like: min, count and max etc.

Operation 5: `Min()` function returns the minimum value in each column.

Operation 6: `Count()` function returns the total number of rows or data samples in the dataset which is 406.

Operation 7: `Max()` function returns the maximum value in each column.

Operation 8: `cars.take(1)` returns the first sample and `cars.take(2)` returns 2 samples.

```
In [137]: cars = sc.textFile("cars.csv")
In [138]: cars
Out[138]: cars has MapPartitionsRDD[213] at textFile at NativeMethodAccessorImpl.java:0

In [139]: cars.show()
Out[139]: +-10.0,0.0,0.207,0.208,0.4334,0,15.0,70.0,1.0+

In [140]: cars.count()
Out[140]: 406

In [141]: cars.max()
Out[141]: "7.0,0.0,389.0,179.0,4390.0,39.9,78.0,1.0"

In [142]: cars.take(1)
Out[142]: ["10.0,0.0,0.207,0.208,0.4334,0,15.0,70.0,1.0"]
```

5. QUESTION 5 –AIRLINE Data:

Airline Data of 2005 from link <http://stat-computing.org/dataexpo/2009/2007.csv.bz2> is downloaded and used for this example.

Operation 1: Read Data

We used 2 ways to read the data. The image shows below how we can import the data inside SPARK using **SPARK CONTEXT** and **SPARK SESSION**.

QUESTION 5

```
In [144]: airlines = sc.textFile("2007.csv")

In [145]: airlines.take(3)

Out[145]: ['Year,Month,DayOfMonth,DayOfWeek,DepTime,CRSDepTime,ArrTime,CRSArrTime,UniqueCarrier,FlightNum,TailNum,ActualElapsedTime,CRSElapsedTime,AirTime,ArrDelay,DepDelay,Origin,Dest,Distance,TaxiIn,TaxiOut,Cancelled,CancellationCode,Diverted,CarrierDelay,WeatherDelay,NASDelay,SecurityDelay,LateAircraftDelay',
'2007,1,1,1,1232,1225,1341,1340,MN,2891,N351,69,75,54,1,7,SMF,ONT,389,4,11,0,,0,0,0,0,0,0',
'2007,1,1,1,1918,1905,2043,2035,MN,462,N370,85,90,74,8,13,SMF,PDX,479,5,6,0,,0,0,0,0,0,0']

In [87]: from pyspark.sql import SparkSession
import pyspark.sql.functions as func
session = SparkSession.builder.appName("AirLines Data").getOrCreate()
dataFrameReader = session.read

In [88]: information = dataFrameReader .option("header", "true") .option("inferSchema", value = True) .csv("2007.csv")
```

Operation 2: Print Header : \$ information.printSchema()

“printSchema()” returns the name of the header/column names of the dataset.

```
In [89]: information.printSchema()

root
|-- Year: integer (nullable = true)
|-- Month: integer (nullable = true)
|-- DayOfMonth: integer (nullable = true)
|-- DayOfWeek: integer (nullable = true)
|-- DepTime: string (nullable = true)
|-- CRSDepTime: integer (nullable = true)
|-- ArrTime: string (nullable = true)
|-- CRSArrTime: integer (nullable = true)
|-- UniqueCarrier: string (nullable = true)
|-- FlightNum: integer (nullable = true)
|-- TailNum: string (nullable = true)
|-- ActualElapsedTime: string (nullable = true)
|-- CRSElapsedTime: string (nullable = true)
|-- AirTime: string (nullable = true)
|-- ArrDelay: string (nullable = true)
|-- DepDelay: string (nullable = true)
|-- Origin: string (nullable = true)
|-- Dest: string (nullable = true)
|-- Distance: integer (nullable = true)
|-- TaxiIn: integer (nullable = true)
|-- TaxiOut: integer (nullable = true)
|-- Cancelled: integer (nullable = true)
-- CancellationCode: string (nullable = true)
-- Diverted: integer (nullable = true)
```

Operation 3: Create Subset : \$

```
airlineSelectedColumns1 = information.select("Month", "UniqueCarrier", "Cancelled",
"LateAircraftDelay") .show()
```

Above command creates a subset of dataset with selected columns using “**.select()**” and showed on console using “**.show()**”. Column names are mentioned as string inside the brackets.

Top 20 results are displayed

```
In [103]: airlineSelectedColumns1 = information.select("Month", "UniqueCarrier", "Cancelled", "LateAircraftDelay")
airlineSelectedColumns1.show()
```

Month	UniqueCarrier	Cancelled	LateAircraftDelay
1	WN	0	0
1	WN	0	0
1	WN	0	31
1	WN	0	3
1	WN	0	0
1	WN	0	0
1	WN	0	0
1	WN	0	1
1	WN	0	0
1	WN	0	24
1	WN	0	0
1	WN	0	0
1	WN	0	3
1	WN	0	42
1	WN	0	0
1	WN	0	0
1	WN	0	0
1	WN	0	30
1	WN	0	0
1	WN	0	0
1	WN	0	0
1	WN	0	9

only showing top 20 rows

Query 1: Apply query on Subset : \$

```
airlineSelectedColumns.groupBy("Month").agg(func.sum("Cancelled")) .orderBy("Month") .show()
```

This commands shows results on selected criteria which in this case is total flights cancelled in each month of year 2007. The data is for year 2007 and the data is print in order of months.

```
In [109]: # airlineSelectedColumns1 .filter(airlineSelectedColumns["LateAircraftDelay"] == "31").show()
airlineSelectedColumns1.groupBy("Month") .agg(func.sum("Cancelled")) .orderBy("Month") .show()
```

Month	sum(Cancelled)
1	15777
2	25465
3	16877
4	11138
5	6841
6	17243
7	13506
8	12295
9	6507
10	7327
11	6279
12	21493

Query 2: Average Delay of Each flight : \$ information.groupBy("FlightNum")

```
.avg("LateAircraftDelay") .show()
```

.avg() is used to calculate the average. Here we have printed average delay of each aircraft. The show is not ordered as “.orderBy()” is not used in this example.

```
In [110]: information.groupBy("FlightNum") .avg("LateAircraftDelay") .show()
```

```
+-----+-----+
|FlightNum|avg(LateAircraftDelay)|
+-----+-----+
|      1959|      1.6505524861878453|
|      2142|      6.8830242510699|
|       496|      4.066509711595056|
|      1342|      8.884453781512605|
|      2659|      5.648972602739726|
|      1238|      6.342790966994789|
|      2122|      1.334035827186512|
|      1829|     11.238139534883722|
|       148|      4.884340320591862|
|      1645|      6.607879924953096|
|      1088|      3.8747815958066396|
|      2866|      5.623124448367167|
|       463|      7.198387096774193|
|       471|      3.2491812227074237|
|      2366|     11.518207282913165|
|      3175|      1.5789473684210527|
|      7340|      6.686182669789227|
|      7253|      6.302222222222222|
|      7240|      2.7661971830985914|
|      5518|      2.9819277108433737|
+-----+-----+
```

only showing top 20 rows

Query 3: Month-wise samples

We can print month-wise samples for each month available in the dataset.

```
In [111]: information.groupBy("Month") .count() .orderBy("Month") .show()
```

```
+-----+-----+
|Month| count|
+-----+-----+
|    1| 621559|
|    2| 565604|
|    3| 639209|
|    4| 614648|
|    5| 631609|
|    6| 629280|
|    7| 648560|
|    8| 653279|
|    9| 600187|
|   10| 629992|
|   11| 605149|
|   12| 614139|
+-----+-----+
```

Query 4: Average and Total distance travelled by each Flight

Command is: `airlineSelectedColumns.groupBy("FlightNum").agg(func.max("Distance"), func.sum("Distance")).orderBy("FlightNum").show()`

First we have created another subset of the main dataset with 3 columns only. Than we apply the query to calculate total distance travelled by a flight and maximum distance travelled by a flight. Top 20 results are also shown below:

```
In [112]: airlineSelectedColumns2 = Information.select("Month", "FlightNum", "Distance")
In [113]: airlineSelectedColumns2.groupBy("FlightNum").agg(func.max("Distance"), func.sum("Distance")).orderBy("FlightNum").show()
```

FlightNum	max(Distance)	sum(Distance)
1	4243	8403820
2	4243	7447691
3	4213	7382791
4	2586	4067920
5	3784	3755377
6	3711	3658765
7	3711	4335473
8	3784	5244639
9	2586	4133897
10	2586	5143218
11	2586	4297478
12	2586	6059674
14	4962	5218246
15	4962	7184979
16	2724	7522454
17	2762	5894324
18	2762	4959386
19	2586	5124826
20	2917	5274152
21	2677	3637952

only showing top 20 rows

This result is displayed in ordered format starting from 1 as we have used “.orderBy()” call.

Query 5: Average and Total distance travelled in each Month

Command is: `airlineSelectedColumns.groupBy("Month").agg(func.max("Distance"), func.sum("Distance")).orderBy("Month").show()`


```
In [114]: airlineSelectedColumns2.groupBy("Month").agg(func.max("Distance"), func.sum("Distance")).orderBy("Month").show()
```

Month	max(Distance)	sum(Distance)
1	4962	441687538
2	4962	402943192
3	4962	458766188
4	4962	441268345
5	4962	453440873
6	4962	456907218
7	4962	474040538
8	4962	473858891
9	4962	431198160
10	4962	450146203
11	4962	433902481
12	4962	446705757

Query 6: Origin of Flights and number of TaxiIn and TaxiOut done

Command 1 is: `airlineSelectedColumns3 = information.select("Origin", "TaxiIn", "TaxiOut")`

Command 2 is: `airlineSelectedColumns3.groupBy("Origin").agg(func.sum("TaxiIn"), func.sum("TaxiOut")).orderBy("Origin").show()`

Here we have created 3rd subset based on origin. The results show sum of flights that “Taxied In” or “Taxied Out” from the “Origin”. Again top 20 results are displayed and order is based on “Origin”.

```
In [121]: airlineSelectedColumns3 = information.select("Origin", "TaxiIn", "TaxiOut")
airlineSelectedColumns3.groupBy("Origin").agg(func.sum("TaxiIn"), func.sum("TaxiOut")).orderBy("Origin").show()
```

Origin	sum(TaxiIn)	sum(TaxiOut)
ABE	48753	76360
ABI	26655	26171
ABQ	272440	432196
ABY	14371	10768
ACK	2898	5483
ACT	27269	20231
ACV	19548	38143
ACY	8051	9800
ADK	374	677
ADQ	2519	4159
AEX	36397	32544
AGS	26688	22522
AKN	855	1117
ALB	103995	217871
ALO	2950	4049
AMA	47584	67263
ANC	100563	273121
APF	7498	5903
ASE	35869	59990
ATL	2448863	8088991

only showing top 20 rows

